I have the following questions related to the manuscript

1. How many modules were actually trained?
2. Which reported performance values are experimentally reproduced?
3. Were all datasets patient independent?
4. How was overfitting prevented?
5. Why is the arthritis model reporting 100% AUC?
6. How is the CardioFusion weighting determined?
7. Why was GPT-2 selected instead of modern medical LLMs?
8. Where is the evaluation of retrieval quality?
9. Why are benchmark values mixed with experimental results?
10. Can the complete source code be released?

I want you people to make the following revisions before submission to any Q1 journal

1. Train every module using real datasets rather than random weights.
2. Remove all benchmark values copied from prior literature unless independently reproduced.
3. Perform external validation on independent datasets.
4. Conduct rigorous statistical significance testing.
5. Add ablation studies for each architectural component.
6. Evaluate retrieval quality using standard RAG metrics.
7. Evaluate generated reports with expert cardiologists.
8. Provide calibration curves and uncertainty estimates.
9. Clearly separate prototype demonstrations from validated experiments.
10. Release source code and trained models for reproducibility.